Content, Not Vocabulary: Evaluating a Style Skill for LLM Prose Against Humanizers, Detectors, and Human Essays

Anonymous ACL submission

Abstract

Advice for making language-model prose read as human usually takes the form of banned-word lists, and "humanizer" tools rewrite text until detectors pass it. We test a different approach: a prompt-level style skill that audits content and structure before vocabulary and forbids adding any fact the writer did not supply. On 16 held-out briefs in ten genres, model judges ranked the skill's writing procedure above plain generation in 47 of 48 judgments, above a word-list prompt in 43 and a popular humanizer's edit of the plain draft in 44, with about a third as many unsupported claims (1.17 against 3.38 per text). We chose it from two frozen candidates after seeing these results; the other also beat every baseline. A judge on a second, same-family model agreed. Two later versions built from judge feedback did not improve on it, and personal-statement rankings shifted with panel composition. Three open-source detectors passed all 89 model-written personal essays we wrote from invented briefs, plain or not, and flagged most pre-ChatGPT admitted college essays; written from notes on real admitted students' essays, model essays failed them too (2 of 24 plain and 11 of 24 skill essays passed, against 1 of 24 real ones). A model counselor rated the skill's coaching of six simulated students above plain feedback, with less invented content. We release the skill, every version, the evaluation data and a rater survey.

1 Introduction

Experienced readers have learned to recognize machine-written prose. Corpus studies trace sharp rises after the release of ChatGPT in words such as *delve* and *underscore* in biomedical abstracts (Kobak et al., 2025) and *commendable* and *meticulous* in peer reviews (Liang et al., 2024), and instruction-tuned models use present participial clauses at two to five times the human rate (Reinhart et al., 2025). Five people who use these tools heavily, voting together, misclassified only 1 of 300 articles, and they cited vocabulary, sentence structure and originality as their evidence (Russell et al., 2025).

Two responses dominate practice. The first is the banned-word list, distributed in prompts, style guides and community "skills". The second is the commercial humanizer, which rewrites AI-generated text to avoid automated or human detection (Roe et al., 2026). Both treat the surface as the problem. Wikipedia's editors, whose guide the most-starred humanizer skill is built on (Chen, 2026), warn against exactly this: "Please do not merely treat these signs as the problems to be fixed; that could just make detection harder" (Wikipedia contributors, 2026). Both responses also ignore a second failure of model prose: models fill gaps with plausible details that nobody supplied (Ji et al., 2023), and in personal writing those details are feelings and memories the writer never reported.

We built and evaluated an alternative: a style skill (a set of instructions an agent loads when a task matches) that orders its audit from content to staging, inflation, rhythm, tone and only then vocabulary, and that never adds a fact absent from the writer's material. For college application essays it switches to a coaching mode and does not write the essay.

Our contributions are:

- An evidence-graded design for a writing skill, built from a review of about 40 sources on AI-writing tells and of 16 existing open-source writing-skill repositories (11 of them humanizer or anti-slop skills), and released with all seven versions.

- A controlled comparison on 20 development and 16 held-out briefs in ten genres against plain generation, a word-list prompt and a popular humanizer, scored by blind judge panels

(including a judge on a second model) and by three open-source detectors.

- A detector audit against human writing collected before ChatGPT. On personal essays the detectors invert: they flag admitted essays and pass model text, plain or not.

- A record of the iterations, including two versions that did not beat their predecessor, and a measurement effect: personal-statement rankings depend on which other texts share the panel.

- A coaching experiment comparing the skill's Coach mode with plain feedback on simulated student drafts.

2 Related work

Stylistic signatures of LLM text. Excess-vocabulary analyses of PubMed and arXiv abstracts and of peer reviews show that LLM adoption shifted the frequency of style words (Kobak et al., 2025; Liang et al., 2024; Juzek and Ward, 2025; Geng and Trotta, 2024), which in PubMed were mostly verbs and adjectives (Kobak et al., 2025). Beyond words, instruction-tuned models write a noun-heavy, informationally dense register (Reinhart et al., 2025), reuse syntactic templates from pre-training data (Shaib et al., 2024), and, in Mandarin, overuse the corrective "not X, but Y" construction, where over half of annotated uses (52.8%) were correct but unneeded (Peng et al., 2026). Professional writers editing model-written creative paragraphs agree on a seven-category taxonomy of idiosyncrasies, led by awkward word choice (28%) and poor sentence structure (20%) ahead of unnecessary exposition and clichés; their detect-then-edit pipeline improves model text but stays below expert edits in expert rankings (Chakrabarty et al., 2025). Our humanizer and Edit conditions are edits of this kind; the skill differs in applying its audit while writing and in forbidding invented facts, and we evaluate it with model judges rather than experts. A study of "slop" built on an expert taxonomy finds that syntactic features predict slop judgments better than lexical ones, though weakly (AUROC 0.64 against 0.49) (Shaib et al., 2026). Banning such patterns token by token makes output unusable at about 2,000 patterns, whereas backtracking at inference time or targeted fine-tuning suppresses thousands of them while keeping quality (Paech et al., 2026). AI generation leaves a stable stylometric footprint that AI editing of human text does not reproduce (Shan et al., 2026).

Detecting AI text. Detectors range from trained classifiers (Adam et al., 2026) to zero-shot statistics (Mitchell et al., 2023; Bao et al., 2024; Hans et al., 2024). The RAID benchmark shows that many are nearly inoperable at low false-positive rates and fail under attacks or unseen generators (Dugan et al., 2024). Paraphrasing defeats watermark-, classifier- and zero-shot detectors (Krishna et al., 2023; Sadasivan et al., 2025), and simple evasion, from adversarial prompting to manual edits and machine paraphrasing, lowers already modest accuracy (Perkins et al., 2024; Weber-Wulff et al., 2023). Detectors misclassify non-native English writing as AI at high rates (Liang et al., 2023) and one RoBERTa detector flagged formal human writing at a 76.3% false-positive rate (Creanga and Dinu, 2026); light, guideline-compliant AI editing is flagged more often than humanizer output (Karr et al., 2026). Untrained humans detect AI text at about chance (Clark et al., 2021; Jakesch et al., 2023), while frequent LLM users do far better (Russell et al., 2025). Text watermarks offer a provenance signal that does not depend on style (Kirchenbauer et al., 2023; Dathathri et al., 2024).

LLMs and admissions essays. LLM-written admissions essays are linguistically distinct from those of 30,000 applicants and resemble the essays of privileged applicants (Lee et al., 2025; Alvero et al., 2024); essay content and style already tracked household income before LLMs (Alvero et al., 2021). Estimated LLM use in essays rose sharply in 2024 (Lee et al., 2026), and in one graduate program, essays flagged as AI-written were associated with lower admission (Isley et al., 2026). AI writing help homogenizes writing across authors (Padmakumar and He, 2024; Doshi and Hauser, 2024; Agarwal et al., 2025) and can shift meaning even when limited to grammar (Abdulhai et al., 2026), and heavy rewriting erases authorship signals (Malik and Awan, 2026).

LLMs as judges. Strong LLM judges agree with human preferences about as often as humans agree with each other, but show position, verbosity and self-enhancement biases (Zheng et al., 2023; Wang et al., 2024), and models can recognize and prefer their own outputs (Panickssery et al., 2024). We

discuss how these biases bear on our design in the Limitations section.

3 The skill

The released skill (version 7) is a Markdown instruction file (173 lines) with four reference files: a catalog of tells with fixes, a personal-essay guide, academic-register rules, and before/after examples. An agent loads it when a writing task matches its description. We built it from two reviews: a survey of 16 open-source writing-skill repositories (11 of them humanizer or anti-slop skills), which supplied the procedure of Chen (2026), the "never inject" guardrails of Bronsdon (2026), and the minimum-edit principle and portability test of Yang (2026) (both MIT-licensed), and an evidence-graded catalog of tells assembled from about 40 sources.

Audit order. The catalog showed that several well-measured signals are structural (participial tails (Reinhart et al., 2025), corrective framing (measured in Mandarin by Peng et al., 2026), importance flags, uniform sentence length), that syntactic features predict expert slop judgments better than lexical ones (Shaib et al., 2026), and that expert readers, although they cite vocabulary most often, also rely on sentence structure and originality (Russell et al., 2025). The skill therefore audits a draft in a fixed order: content (the "portability test": could this sentence move unchanged into someone else's text?), staging, inflation, stated lessons, rhythm, paragraph endings, tone, vocabulary, and format. A strong tell justifies an edit on one sighting; weak tells, such as a single transition word, only when several cluster.

Hard rules. The skill never adds a fact, name, number, quote, feeling or memory that the user or the source did not provide; where a sentence needs one, it asks or leaves a bracketed question. It keeps every claim and hedge when editing, uses no evasion tricks (typos, rare synonyms, look-alike characters), and makes minimal edits to human-written text.

Over-correction checks. Humanizer rewrites create tells of their own: a blind test by a contributor, linked from the humanizer's README, counted more uniform-rhythm tells (60 to 101) and more aphoristic closers (38 to 61) after the rewrite pass (Lenk, 2026). The final check lists the last sentence of every paragraph and rewrites when more than two are short or quotable, and checks that a personal essay still has an arc and a line of reflection.

Modes. *Write* runs the full procedure; for personal essays it drafts in two passes, first for shape and then as an editor of its own draft (Section 5.4). *Edit* applies the audit to the user's own draft with minimal changes. *Rebuild*, for drafts the user says an AI wrote, first checks every fact, quote and reading against the user's material; if the draft was built from that material it keeps the checked evidence and writes fresh sentences, and if it was built from the topic alone it treats the draft's examples and plan as the stock answer and writes from the first step. *Detect* quotes passages and names patterns without guessing authorship. *Coach*, used for application essays, interviews the student, proposes a structure, critiques drafts and lists corrections, and does not write the essay or example sentences built from the student's story.

4 Evaluation

We evaluated the skill in three rounds, after a seven-brief pilot that shaped versions 1–4 (Table 2): a development round on 20 briefs, a version comparison on the same briefs, and a held-out round on 16 new briefs written after versions 4 and 6 were frozen. The released version 7 was assembled after the held-out round. A separate experiment tests Coach mode.

Briefs. The development set has 20 briefs in five genres, four each (Appendix A, Table 3). Four fictional students, each with about 1,100 to 1,600 words of interview notes and answers to the skill's depth questions, supply the material for a Common App personal statement and a 250-word supplemental answer. The other briefs are 11th-grade argument essays and science-club explainers with no facts file, and literary analyses restricted to four passages from a public-domain novel or play, copied from Project Gutenberg and checked by script. The held-out set has 16 briefs: two new students (a personal statement and a short answer each), two arguments, two explainers, two literary analyses, and six briefs in five genres outside the development set: two emails (an extension request to a teacher and a cold email to a professor), a cover letter, a team blog post, the discussion section of a lab report, and a school news article. The seven-brief pilot that shaped versions 1–4 had included a first-person blog post and a research-report meth-

ods section. Every brief names the only permitted source of specifics; the people in the facts files are fictional.

Conditions. Fresh [generator model, name and version] agents wrote every text. **Plain** (P): the model's default writing. **Word list** (L): the same request plus an instruction to avoid 24 common AI words, six stock phrases and constructions, and em dashes and to vary sentence length. **Humanizer edit** (H): the most-starred of the 11 humanizer skills we surveyed (about 59k GitHub stars on 20 October 2026) (Chen, 2026), applied to P's draft. It edits a draft rather than writing from the brief, so its like-for-like comparison is with the skill's Edit or Rebuild mode on the same draft, not with Write mode. The skill conditions were the Write mode, the Edit mode applied to P's draft, and from version 5 on a Rebuild mode for drafts the user says an AI wrote. Every prompt carried the same brief and facts-file instruction; on briefs about fictional people, the skill and humanizer prompts also said the people were fictional, so that the skill would write the piece instead of switching to Coach mode. The development round compared P, L, H, Edit and Write (version 4); the version comparison compared P, L, version 4 Write, version 5 Write and version 5 Rebuild; the held-out round compared P, L, H, version 4 Write, version 6 Write and version 6 Rebuild.

Judges. Judge agents read all texts for a brief in a shuffled order under neutral labels (T1, T2, . . .), but each judge opened the texts by file path, and the path contained the condition's folder name (a single letter such as P or H; D0, RS, RP, FS or FP in the coaching experiment). Blinding was therefore imperfect: in 6 of the 128 writing-panel judgments the explanation names each text's condition letter (for example "T3 (L)"), and the coaching counselor labelled every piece of feedback by its path. The judges were a strict editor, and the genre's real reader (an admissions officer, an English or science teacher, the recipient of the email, a hiring manager, a team mentor, a newspaper adviser). Each ranked the texts, scored each from 1 to 10 on quality, human-likeness, specificity, voice and structure, rated AI suspicion from 1 to 5, and quoted every claim not supported by the facts file (or every invented statistic, source or factual error when there was none), a coarser, judge-based version of checking generated facts against a knowledge source (Min et al., 2023). In the held-out round a third judge, a second genre reader, ran on a different model of the same family ([second judge model, name and version]). We report mean rank with bootstrap confidence intervals over judgments, a Friedman test across conditions on per-brief mean ranks, and pairwise Wilcoxon signed-rank tests on per-brief mean ranks, so the unit of every test is the brief (16 held-out), not the judgment (48). Judgments of one brief rate the same texts, and the two same-model judges agreed closely (mean Kendall τ 0.78), so the intervals understate uncertainty: resampling briefs instead widens version 4's held-out interval from [1.62, 2.23] to [1.52, 2.38]. Over the 15 held-out pairwise tests, Holm's correction keeps every comparison we report as significant below an adjusted p of 0.025.

Detectors. We scored the development-round and held-out texts with three open-source classifiers: D1, a DeBERTa-v3-large model trained on RAID (Desklib, 2025); D2, a RoBERTa-base model (Fakespot, 2025); and D3, a RoBERTa-large model trained on a RAID subset plus generations that include Anthropic-model output (SuperAnnotate, 2024). Texts are split into chunks of about 350 words, and document scores are length-weighted means. We report the share below 20%, the threshold below which Turnitin shows only an asterisk instead of an AI score (Turnitin, 2026).

Human baselines. To see how the detectors treat human writing, we scored three sets collected before ChatGPT: 16 admitted college essays published by Hamilton College (classes of 2018 and 2022; Hamilton College Office of Admission, 2026) and Connecticut College (class of 2025; Connecticut College Office of Admission, 2026); 18 high-scoring argument essays from the PERSUADE 2.0 corpus (Crossley et al., 2024) (15 with the top score of 6, three with 5); and three excerpts from two science explainers published in 2020 (one by Bert Hubert, two from one article in The Conversation).

Coaching experiment. For each of the six fictional students, an agent playing the student wrote a first draft of the personal statement from their notes, with no style guidance. The draft then received feedback from the skill's Coach mode (version 6 files) or from the plain model asked for feedback, and the student agent revised it once with each set of feedback. Two blind judges ranked the first draft and both revisions, and a counselor agent rated both

pieces of feedback for usefulness and counted the essay-ready sentences each wrote for the student. We also measured how much of each revision was copied from the feedback it received: the share of the revision's word 6-grams that occur in the feedback but in neither the first draft nor the student's notes.

5 Results

5.1 Held-out briefs

Table 1 and Figure 1 give the main result. The 16 held-out briefs, in ten genres, were written after versions 4 and 6 were frozen, but we chose which of the two to release after seeing these results, so version 4's figures are the better of two candidates and somewhat optimistic for the released skill; version 6, the other candidate, also beat all three baselines (Table 1). Version 4's Write mode ranked first in 23 of 48 judgments and had a mean rank of 1.92 out of 6, against 5.65 for plain generation (Friedman $\chi^2 = 51.0$, $p < 10^{-8}$). It ranked above plain generation in 47 of 48 judgments, above the word-list prompt in 43 and above the humanizer edit in 44 (Wilcoxon $p \leq 0.003$ for each). It made 1.17 unsupported claims per text, against 3.38 for plain generation and 2.98 for the word-list prompt. The Rebuild mode, applied to the plain draft, ranked above the humanizer applied to the same draft in 42 of 48 judgments ($p = 0.0008$) and had the fewest unsupported claims of any condition (0.79).

Version 6, which added rules derived from a diagnosis of the development judgments, did not improve on version 4: it ranked above version 4 in 21 of 48 judgments ($p = 0.35$). Its argument essays were built around a worked numerical case; one judge called its Write-mode essay on graded homework "dry and mechanical," and another called its Rebuild of the same brief "arithmetic-heavy for an English class." We therefore released version 7: version 4's Write procedure, with six examples that came from development briefs replaced, plus version 6's Rebuild mode, adapted so that drafts built from the topic alone go to version 4's Write procedure (Section 3). Version 7 as a whole did not go through the held-out panels; its only direct check is a smoke test on three held-out plain drafts, so the held-out numbers describe its parts.

The third judge, which ran on a different model, gave the same order (version 4 1.94, version 6 2.25, Rebuild 2.44, word list 4.19, humanizer 4.81, plain 5.38), so within one model family the ranking does not depend on the judge model.

Condition	Rank	1st	Human	Unsup.
Plain	5.65	0	4.48	3.38
Word list	4.17	3	6.48	2.98
Humanizer edit	4.25	2	5.69	1.88
Skill Write (v4)	**1.92**	**23**	**7.15**	1.17
Skill Write (v6)	2.42	17	6.96	0.96
Skill Rebuild (v6)	2.60	3	7.10	**0.79**

Table 1: Held-out evaluation: 16 briefs, three blind judges each (48 judgments). Rank: mean rank among six texts (1 = best); 1st: first places; Human: judged human-likeness (1–10); Unsup.: claims per text that judges found unsupported by the facts file or wrong.

By genre, version 4 ranked best on personal statements (1.33), arguments (1.33), literary analysis (1.33), the emails (1.67) and the news article (1.00). It ranked worst among the skill conditions on the blog post (4.67), where the word-list prompt won (1.67), and the word-list prompt also edged it on the cover letter (2.00 against 2.33). With one or two briefs per genre these are descriptive. Both losses are in first-person genres, where the audit may strip the looseness the genre wants, but each rests on one brief and three judgments, and the cover-letter gap (2.00 against 2.33) is small. The no-invention rule also did little where there was no source file: version 4's argument essays had 3.17 unsupported or wrong claims per text against 4.00 for plain generation, and its explainers 1.50 against 1.33, so most of the overall drop comes from the genres with a facts or quotes file.

5.2 Development round and the version comparison

On the 20 development briefs, version 4's Write mode ranked first in 28 of 40 judgments (mean rank 1.45 of 5; plain 4.62; word list 2.27; humanizer 3.48; Edit mode on the plain draft 3.17). Version 5 added a plain student voice for school essays, a Rebuild mode and a personal-essay checklist. In a new five-way panel on the same briefs it raised judged human-likeness (7.45 against 6.90) but ranked below version 4 in 28 of 40 judgments (Wilcoxon $p = 0.055$, so it did not improve on version 4 rather than clearly doing worse): judges caught simplified reasoning, and the example stance line in the rule ("I think the policy is wrong") became the opening of three of four argument essays.

The version comparison also exposed a measurement problem. The plain and version 4 texts were identical in both panels, yet on personal statements

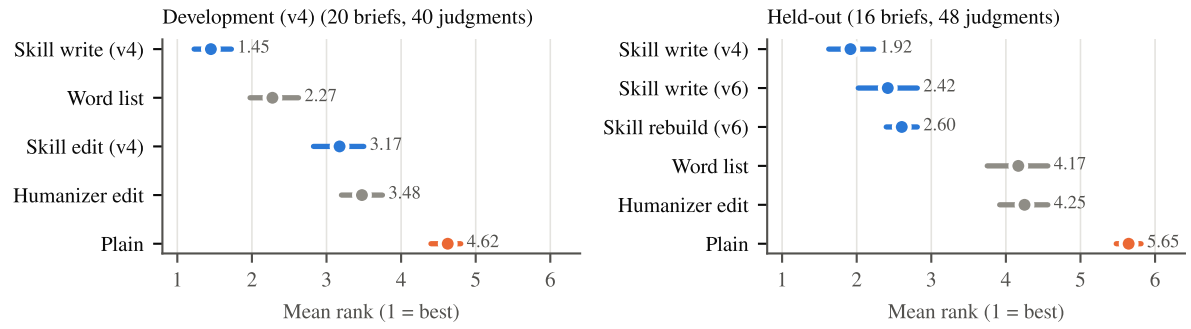


Figure 1: Mean rank with 95% bootstrap intervals (1 = best). Left: development briefs, five texts per panel. Right: held-out briefs, six texts per panel, three judges each. Blue: the skill; orange: plain generation; grey: other approaches.

version 4 ranked above plain in 8 of 8 judgments in the first panel and in only 3 of 8 in the second. The first panel contained two edits of the plain draft; judges noticed the shared sentences and may have marked the whole family down. On arguments, explainers and literary analyses the comparison held (8 of 8 in both panels); on short answers it went from 7 of 8 to 6 of 8. Rankings of personal statements depend on what they are ranked against, and the held-out panels do not escape this: they again include the humanizer's edit of the plain draft, which shares 44% of its word 6-grams with that draft (48% and 55% for the two edits in the first development panel, 18% for the rebuild in the second), so version 4's held-out lead over plain generation on personal statements and short answers (12 of 12 judgments) may be inflated in the same way.

5.3 Detectors and human writing

Figure 2 shows the held-out scores. All 16 texts written with version 4's Write mode and all 16 version 6 rebuilds scored below 20% on all three detectors (version 6's Write mode: 15 of 16), against 9 of 16 for plain generation and 12 each for the word-list prompt and the humanizer edit. The development round agreed (17 of 20 Write-mode texts against 9 of 20 plain; the Edit mode passed 10 of 20). The advantage came from the school and technical genres: on personal statements and short answers every condition passed, plain generation included, with no score above 3% on the held-out briefs and none above 6.2% in the development round. The human baselines went the other way. None of the 16 admitted college essays scored below 20% on D3, whose mean for them was 61%, and only 5 scored below 20% on both D1 and D2; 2 of the 18 PERSUADE essays passed all three. On personal essays written from our invented briefs these detectors passed every model-written text (89 of 89 across all rounds) and flagged most human ones. That pass rate did not hold for real material: when each model essay was written from notes reverse-engineered from a real admitted essay (24 development essays of the human-anchored benchmark), all three detectors scored below 20% for only 2 of 24 plain essays and 11 of 24 skill essays, against 1 of the 24 real ones.

5.4 Personal statements: arc against fidelity

Personal statements were the hardest genre throughout development (Table 2). Plain generation built a better arc by writing connective feelings and causes the student never stated (3 to 11 unsupported claims per essay), and judges preferred it. Forbidding invention flattened the skill's first three versions; one judge wrote that a version 3 essay "mostly walks through the notes in order" and that its climax "lands in paragraph 7 of 9." Two-pass drafting in version 4 (write for shape first, then audit the draft as an editor) recovered part of the arc: the Write mode's first personal-statement win against plain generation on the original notes (1 of 2 judgments), with 4 unsupported claims against 14. Giving both conditions the student's answers to depth questions raised everyone's quality and halved plain generation's invention. On the two held-out students, version 4 ranked first in 5 of 6 personal-statement judgments with 0.67 unsupported claims per essay, against 4.83 for plain generation.

5.5 Coaching

Both kinds of feedback improved the first drafts: the draft ranked last in all 12 judgments. The revision made with Coach-mode feedback ranked

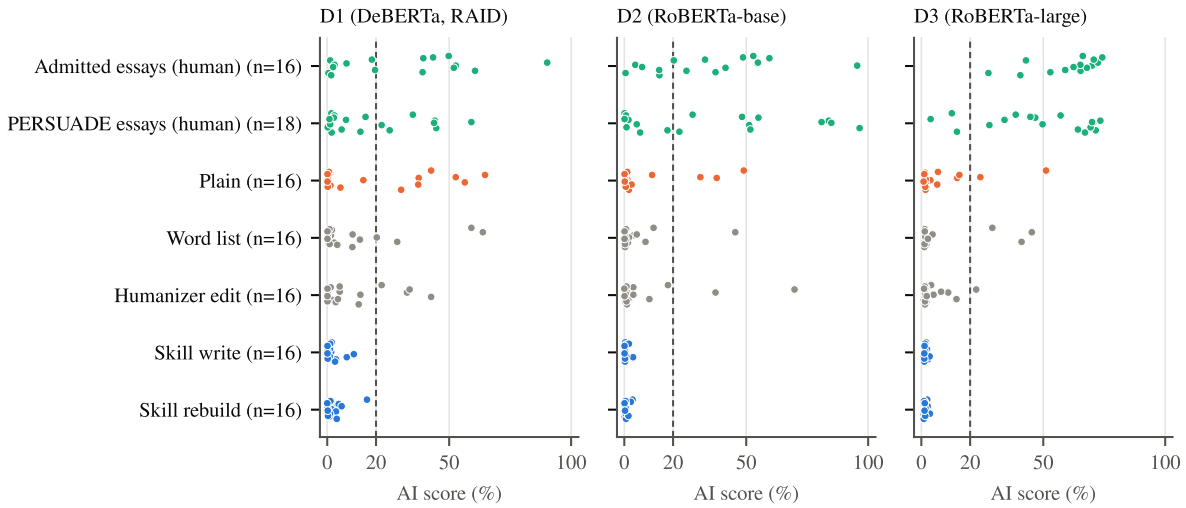


Figure 2: Detector scores per text on the held-out briefs (Skill write: version 4; Skill rebuild: version 6), with human writing collected before ChatGPT for comparison (green). D1–D3 are the detectors described in Section 4. The dashed line is 20%. D3's scores saturate near 73%.

Version	Change	vs. plain	Unsup.
v1	audit while drafting	0–4	–
v2	stricter sample-essay rules	0–4	–
v3	countable checks	0–4	1 vs 16
v4	two-pass drafting	1–3	6 vs 19

Table 2: Personal-statement iterations in development (pairwise judgments of better writing, skill wins–plain wins; depth notes for v3–v4; unsupported claims summed over the same pairs, skill vs plain). Small samples: four judgments (two text pairs, two judges) per row. v1–v2 used the original notes, where unsupported claims were 16 vs 29 and 6 vs 33.

above the one made with plain feedback in 7 of 12 judgments (not significant), with fewer unsupported claims (1.67 against 2.75) and lower AI suspicion (1.83 against 2.08). The clearer difference was in the feedback itself. The counselor judge preferred Coach mode's feedback for all six students (usefulness 8.67 against 7.17 of 10) and found less invented content in it (1.17 items per student against 5.17). Neither kind of feedback wrote many whole sentences for the student (1.33 per response for plain feedback; 0.67 for Coach mode, whose rules forbid any, with three in one response), but plain feedback supplied the essay's interpretation: one response supplied what the counselor called "the essay's interpretive center," a line the student never said, and framed it as his own answer; the counselor counted this as substantive content of the kind admissions policies forbid (Common App, 2023; Yale Undergraduate Admissions, 2026).

The simulated students copied almost nothing from either feedback (on average under 1% of word 6-grams; at most 3.7% in one revision), so this experiment cannot measure what a rushed real student would copy. Coach mode's responses were also about 1.5 times longer (2,263 words against 1,468), which may be more than a student will read and may have helped them with the counselor, since model judges tend to favor longer answers (Zheng et al., 2023). The feedback comparison rests on one counselor judgment per student, six in all, from a model that was told the Common App rule Coach mode cites and that could see each condition's folder name in the file path. The experiment ran version 6's personal-essay reference, whose Coach-mode section differs from the released version 7's, and four of the six students are development personas.

6 Discussion

Why content first. Judges' reasons point to the same mechanism across genres. The texts they ranked first argued from one specific organizing idea, used the writer's own facts and kept them straight, and ended on content; the texts they ranked last had the stock plan of the genre (a three-reason thesis, a sensory hook, a summary close). The word-list prompt changed the vocabulary and the voice, and judges rated it more human than plain generation, but it kept the stock plan and invented almost as much (2.98 unsupported claims per text against 3.38). The humanizer edit removed tells

sentence by sentence and kept the draft's plan, and judges noticed when two texts were edits of the same draft. We read the skill's gains as coming from what a vocabulary pass cannot reach (what the text claims, in what order, and from which facts), but that reading rests on judges' comments. We did not ablate the audit order, the no-invention rule or two-pass drafting, and the skill gives the model far longer instructions than the word-list prompt, so the evaluation shows that the skill as a whole helps, not which part does.

Iterating on judges. Diagnosing judge comments produced plausible rules that did not transfer. Version 5's student-voice rule raised human-likeness and lowered reasoning quality, and its example line became a template; version 6's argument rules, distilled from 16 judgments that praised one move, produced held-out argument essays that judges called dry, mechanical and, in one case, arithmetic-heavy. Both regressions were caught only because each version was compared with version 4 in the same panel: version 5 on the development briefs and version 6 on the held-out briefs. We think style skills, like models, need a held-out set: the development judges' taste is part of what gets fitted.

Detectors in admissions. On personal essays the three detectors got the direction wrong: they scored real admitted essays high and model essays low, with or without the skill. A plausible reason is domain. Two of the detectors were trained on RAID, whose eight domains (scientific abstracts, book summaries, news, poetry, recipes, Reddit posts, reviews and Wikipedia) do not include the personal essay (Dugan et al., 2024). In most of those domains human writing is uneven and specific and model writing is polished and general. Admitted essays are polished, edited many times, and full of the reflective phrasing that admissions advice recommends, so they look like the model side of the training data; a model told to use concrete, uneven detail looks like the human side. Whatever the cause, a score from these three detectors is not evidence that an applicant used AI; commercial detectors may behave differently, and we did not test them. The Common App and colleges that forbid substantive AI content (Common App, 2023; Brown University Office of College Admission, 2026; Yale Undergraduate Admissions, 2026) have to enforce it some other way, and our results, together with evidence that detector scores rise with formal, academic-register writing (Creanga and Dinu, 2026; Karr et al., 2026), suggest that style-based detection would fall hardest on applicants who write in the polished register that essay coaching teaches.

Coaching versus ghostwriting. In development, personal statements were the one genre where the skill's writing did not clearly win (on the two held-out students it ranked first, in panels that may favor it; Section 5). The essays that judges liked most built an arc, a turning point and a closing insight, and plain generation built them by inventing feelings and causes the notes did not contain. When the skill forbade that, it lost some of the arc. The material that makes an application essay work (what the student actually thought, what they still have not resolved) exists only in the student's head, so the honest way for a model to help is to ask for it. Coach mode does this and leaves the writing to the student, so the essay's text stays the student's own. Policies differ on the rest: Yale permits grammar and spelling review and "general advice or topic suggestions at the start of the writing process" (Yale Undergraduate Admissions, 2026), while Brown permits only spelling and grammar review (Brown University Office of College Admission, 2026), so whether a proposed structure and draft critique are allowed depends on the school.

7 Conclusion

A style skill that fixes content and structure before vocabulary, and that never adds facts, produced writing that model judges from one model family preferred to plain generation, a word-list prompt and a popular humanizer on briefs it was not developed on (though we chose it from two candidates on those briefs), with about a third of plain generation's unsupported claims. A model counselor rated its Coach-mode feedback more useful than plain feedback, and that feedback left the writing to the student. The detectors we tested scored its personal essays, and plain model essays, as human and flagged most real admitted essays, so they should not be used as evidence in admissions. The open questions are whether human readers agree with the model judges, which our released survey is built to test, and how to keep casual first-person genres loose while auditing them.

Limitations

Model judges. All judges are language models from the same family as the generator, and models can recognize and favor their own output (Panickssery et al., 2024). Self-preference does not by itself favor one of our conditions, since every condition was produced by the same model, but judges may share the generator's taste for a particular kind of polish. We replaced condition names with neutral labels (though file paths still showed them; Section 4), shuffled the order of texts, used two judge roles (editor and genre reader), and in the held-out round added a judge running on another model of the same family; we did not test judges from other model families or recruit enough human raters. The circularity goes beyond self-preference: agents from the same family drafted the skill, versions 5 and 6 were built from these judges' comments, and the editor prompt in every round, like most reader prompts, cast the judge as someone watching for AI writing ("knows the research on AI-writing tells", "wary of AI-written work"), which is what the skill was built to remove. The held-out round is held out over briefs, not over judges, so our results show what judges from this family prefer. The two same-model judges also agreed with each other (mean Kendall τ 0.78) more than with the third judge (0.55–0.62). A human rating study with the released survey is the most important missing piece.

Test material. The student personas are fictional and were written by a model; real interview notes are likely to be shorter, messier and less quotable. Coach mode was checked on one real student draft and one pressure test, which shows it can work but not how often it does.

Detectors. We used three open-source classifiers. Commercial detectors used by schools (Turnitin, GPTZero) were not tested, so the pass rates here do not predict their scores. The human baselines are small (16 admitted essays, 18 PERSUADE essays) and the admitted essays were selected and edited for publication by the colleges.

Scope and development. One generator model wrote every text, in English, for genres common in US high schools. Versions 1–4 were developed on a seven-brief pilot (Table 2), and five of those briefs (one student's two essays, the phone-ban argument, the mRNA explainer and the Gatsby analysis) are in the 20-brief development set; version 4 also contained six example phrases taken from development briefs. Version 4 ranked about the same on those five briefs (mean 1.40) as on the other fifteen (1.47). Version 5 was built from the development-round judgments and version 6 from both development panels, all on the same 20 briefs, so version 5's version-comparison scores are in-sample. The held-out round uses 16 new briefs, but they were written by the same team that built the skill. Per-genre results rest on one to four briefs each (one brief in each of four held-out genres) and should be read as descriptive; each condition produced one text per brief, so a result on a brief also reflects the luck of that single sample.

Ethical considerations

Dual use. The skill lowers detector scores as a side effect of better writing, and a student could use it to pass off model-written work as their own. We took three steps. The skill uses no evasion tricks (typos, homoglyphs, invisible characters) of the kind studied in the attack literature (Krishna et al., 2023; Sadasivan et al., 2025), but it lowers detector scores anyway: on the six held-out argument, explainer and literary-analysis briefs, every plain text was flagged by at least one detector and every text version 4 wrote passed all three. Its integrity rule sends work a student must submit as their own to Coach mode only when the user says so; our evaluation prompts asked for Write mode on 11th-grade assignments and got finished essays, and a student could do the same with homework. For application essays it coaches and does not write, and it tells the student that the Common App, Yale and many programs count substantive AI-written content as not the applicant's own work. And it tells users that recent model output may carry a provenance watermark (Anthropic, 2026; BleepingComputer, 2026) and that the skill is not a way to hide AI use.

Detector developers. Because we report low detector scores for model text and high scores for human essays, we [notified the developers of all three detectors on <date>], at least 30 days before any preprint.

Human writing and raters. We scored published admitted essays and PERSUADE essays and report only their scores; we do not redistribute their texts, and the released score files identify each essay by source and number rather than by its writer's

name, so no named student is linked to an AI score. No real applicant's essay appears in the released data. Any human rating study uses volunteers who give informed consent, collects no names or contact details, and [had approval from <school or teacher>]. Raters under 18 take part only with a parent's or guardian's permission, which the survey asks for but cannot verify. The generated texts and the survey's answer key are in our public repository, so a rater who searched for a passage could learn its condition; the survey asks raters not to.

AI use statement

AI agents ([generator model, name and version], Anthropic) did a large share of this work: they surveyed the literature and existing skills, drafted every version of the skill, generated all evaluation texts, acted as judges, wrote the analysis scripts, and drafted this paper. The development log, prompts, scripts and every intermediate version are released. [The author describes here, accurately, what they did themselves: e.g. which decisions they made, which numbers and citations they checked against the data and sources, and which sections they rewrote.]

References

Marwa Abdulhai, Isadora White, Yanming Wan, Ibrahim Qureshi, Joel Z. Leibo, Max Kleiman-Weiner, and Natasha Jaques. 2026. How LLMs distort our written language. *arXiv preprint arXiv:2603.18161*.

George Alexandru Adam, Alexander Cui, Edwin Thomas, Emily Napier, Nazar Shmatko, Jacob Schnell, Jacob Junqi Tian, Alekhya Dronavalli, Edward Tian, and Dongwon Lee. 2026. GPTZero: Robust detection of LLM-generated texts. *arXiv preprint arXiv:2602.13042*.

Dhruv Agarwal, Mor Naaman, and Aditya Vashistha. 2025. AI suggestions homogenize writing toward Western styles and diminish cultural nuances. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, pages 1–21. Association for Computing Machinery.

A. J. Alvero, Jinsook Lee, Alejandra Regla-Vargas, René F. Kizilcec, Thorsten Joachims, and Anthony Lising Antonio. 2024. Large language models, social demography, and hegemony: comparing authorship in human and synthetic text. *Journal of Big Data*, 11(1):138.

AJ Alvero, Sonia Giebel, Ben Gebre-Medhin, Anthony Lising Antonio, Mitchell L. Stevens, and Benjamin W. Domingue. 2021. Essay content and style are strongly related to household income and SAT scores: Evidence from 60,000 undergraduate applications. *Science Advances*, 7(42):eabi9031.

Anthropic. 2026. How Claude marks AI-generated content. Claude Help Center, https://support.claude.com/en/articles/16266773-how-claude-marks-ai-generated-content. Accessed 20 October 2026.

Guangsheng Bao, Yanbin Zhao, Zhiyang Teng, Linyi Yang, and Yue Zhang. 2024. Fast-DetectGPT: Efficient zero-shot detection of machine-generated text via conditional probability curvature. In *The Twelfth International Conference on Learning Representations*.

BleepingComputer. 2026. How Anthropic plans to watermark Claude's AI-generated text. https://bleepingcomputer.com/news/artificial-intelligence/how-anthropic-plans-to-watermark-claudes-ai-generated-text. News report, accessed 20 October 2026.

Conor Bronsdon. 2026. avoid-ai-writing. GitHub repository, https://github.com/conorbronsdon/avoid-ai-writing. Agent skill, version 3.38.0; MIT license; accessed 20 October 2026.

Brown University Office of College Admission. 2026. Integrity in the application process. https://admission.brown.edu/apply/how-apply/integrity-application-process. Accessed 20 October 2026.

Tuhin Chakrabarty, Philippe Laban, and Chien-Sheng Wu. 2025. Can AI writing be salvaged? mitigating idiosyncrasies and improving human-AI alignment in the writing process through edits. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, pages 1–33. Association for Computing Machinery.

Siqi Chen. 2026. Humanizer. GitHub repository, https://github.com/blader/humanizer. Agent skill (SKILL.md metadata version 3.1.0); MIT license, copyright 2025 Siqi Chen; accessed 2026-10-20.

Elizabeth Clark, Tal August, Sofia Serrano, Nikita Haduong, Suchin Gururangan, and Noah A. Smith. 2021. All that's 'human' is not gold: Evaluating human evaluation of generated text. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7282–7296, Online. Association for Computational Linguistics.

Common App. 2023. Fraud policy. https://www.commonapp.org/fraud-policy/. Last updated August 2023; accessed 20 October 2026.

Connecticut College Office of Admission. 2026. Essays that worked. https://www.conncoll.edu/a

dmission/apply/essays-that-worked/. Essays by the class of 2025.

Claudiu Creanga and Liviu Dinu. 2026. Ontological instability and statistical amplification: The paradox of "humanizing" LLM-generated text. *arXiv preprint arXiv:2610.03110*.

S. A. Crossley, Y. Tian, P. Baffour, A. Franklin, M. Benner, and U. Boser. 2024. A large-scale corpus for assessing written argumentation: PERSUADE 2.0. *Assessing Writing*, 61:100865.

Sumanth Dathathri, Abigail See, Sumedh Ghaisas, Po-Sen Huang, Rob McAdam, Johannes Welbl, Vandana Bachani, Alex Kaskasoli, Robert Stanforth, Tatiana Matejovicova, Jamie Hayes, Nidhi Vyas, Majd Al Merey, Jonah Brown-Cohen, Rudy Bunel, Borja Balle, Taylan Cemgil, Zahra Ahmed, Kitty Stacpoole, and 5 others. 2024. Scalable watermarking for identifying large language model outputs. *Nature*, 634(8035):818–823.

Desklib. 2025. desklib/ai-text-detector-v1.01. Hugging Face model card, https://huggingface.co/desklib/ai-text-detector-v1.01. DeBERTa-v3-large fine-tuned on RAID; MIT license; last updated 17 Feb 2025; accessed 2026-10-20.

Anil R. Doshi and Oliver P. Hauser. 2024. Generative AI enhances individual creativity but reduces the collective diversity of novel content. *Science Advances*, 10(28):eadn5290.

Liam Dugan, Alyssa Hwang, Filip Trhlík, Andrew Zhu, Josh Magnus Ludan, Hainiu Xu, Daphne Ippolito, and Chris Callison-Burch. 2024. RAID: A shared benchmark for robust evaluation of machine-generated text detectors. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12463–12492, Bangkok, Thailand. Association for Computational Linguistics.

Fakespot. 2025. fakespot-ai/roberta-base-ai-text-detection-v1. Hugging Face model card, https://huggingface.co/fakespot-ai/roberta-base-ai-text-detection-v1. RoBERTa-base fine-tune; Apache-2.0 license; last updated 25 Feb 2025; accessed 2026-10-20.

Mingmeng Geng and Roberto Trotta. 2024. Is ChatGPT transforming academics' writing style? *arXiv preprint arXiv:2404.08627*.

Hamilton College Office of Admission. 2026. College essays that worked. https://www.hamilton.edu/admission/apply/college-essays-that-worked. Essays by the classes of 2018 and 2022.

Abhimanyu Hans, Avi Schwarzschild, Valeriia Cherepanova, Hamid Kazemi, Aniruddha Saha, Micah Goldblum, Jonas Geiping, and Tom Goldstein. 2024. Spotting LLMs with Binoculars: Zero-shot detection of machine-generated text. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 17519–17537. PMLR.

Calvin Isley, Johann D. Gaebler, and Sharad Goel. 2026. AI-written admissions essays are widespread but penalized. *arXiv preprint arXiv:2609.22549*.

Maurice Jakesch, Jeffrey T. Hancock, and Mor Naaman. 2023. Human heuristics for AI-generated language are flawed. *Proceedings of the National Academy of Sciences*, 120(11):e2208839120.

Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38.

Tom S. Juzek and Zina B. Ward. 2025. Why does ChatGPT "delve" so much? exploring the sources of lexical overrepresentation in large language models. In *Proceedings of the 31st International Conference on Computational Linguistics*, pages 6397–6411, Abu Dhabi, UAE. Association for Computational Linguistics.

Jonathan A. Karr, Jr., Grigorii Khvatskii, Ting Hua, and Nitesh V. Chawla. 2026. Why AI detection fails for academic integrity. *arXiv preprint arXiv:2608.11256*. ArXiv comment: accepted to ACM AI Leadership Summit.

John Kirchenbauer, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. 2023. A watermark for large language models. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 17061–17084. PMLR.

Dmitry Kobak, Rita González-Márquez, Emőke-Ágnes Horvát, and Jan Lause. 2025. Delving into LLM-assisted writing in biomedical publications through excess vocabulary. *Science Advances*, 11(27):eadt3813.

Kalpesh Krishna, Yixiao Song, Marzena Karpinska, John Wieting, and Mohit Iyyer. 2023. Paraphrasing evades detectors of AI-generated text, but retrieval is an effective defense. In *Advances in Neural Information Processing Systems*, volume 36, pages 27469–27500.

Jinsook Lee, AJ Alvero, Thorsten Joachims, and René Kizilcec. 2025. Poor alignment and steerability of large language models: Evidence from college admission essays. *arXiv preprint arXiv:2503.20062*.

Jinsook Lee, Conrad Borchers, AJ Alvero, Thorsten Joachims, and Rene F. Kizilcec. 2026. The digital divide in generative AI: Evidence from large language model use in college admissions essays. In *Proceedings of the Thirteenth ACM Conference on Learning @ Scale*, pages 179–190. Association for Computing Machinery.

Ryan Lenk. 2026. Blind study: the rewrite pass wins 16/16 on quality, and does not move AI detection. GitHub issue #229, blader/humanizer, https://github.com/blader/humanizer/issues/229. Opened 19 August 2026; accessed 21 October 2026.

Weixin Liang, Zachary Izzo, Yaohui Zhang, Haley Lepp, Hancheng Cao, Xuandong Zhao, Lingjiao Chen, Haotian Ye, Sheng Liu, Zhi Huang, Daniel A. McFarland, and James Y. Zou. 2024. Monitoring AI-modified content at scale: A case study on the impact of ChatGPT on AI conference peer reviews. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 29575–29620. PMLR.

Weixin Liang, Mert Yuksekgonul, Yining Mao, Eric Wu, and James Zou. 2023. GPT detectors are biased against non-native English writers. *Patterns*, 4(7):100779.

Ushna Malik and Moiz Sadiq Awan. 2026. The assistant erased you: Measuring loss of authorship signals in AI-mediated communication. *arXiv preprint arXiv:2608.00926*.

Sewon Min, Kalpesh Krishna, Xinxi Lyu, Mike Lewis, Wen-tau Yih, Pang Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12076–12100, Singapore. Association for Computational Linguistics.

Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D. Manning, and Chelsea Finn. 2023. DetectGPT: Zero-shot machine-generated text detection using probability curvature. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 24950–24962. PMLR.

Vishakh Padmakumar and He He. 2024. Does writing with language models reduce content diversity? In *The Twelfth International Conference on Learning Representations*.

Samuel J. Paech, Allen G. Roush, Judah Goldfeder, and Ravid Shwartz-Ziv. 2026. Antislop: A comprehensive framework for identifying and eliminating repetitive patterns in language models. In *The Fourteenth International Conference on Learning Representations*. ArXiv:2510.15061.

Arjun Panickssery, Samuel R. Bowman, and Shi Feng. 2024. LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems*, volume 37, pages 68772–68802.

Yihang Peng, Yi Jing, Dewen Liu, Weihan Meng, and Ying Liu. 2026. Not X, but when: Diagnosing a Mandarin LLM writing tic through contextual appropriateness. *arXiv preprint arXiv:2610.22249*.

Mike Perkins, Jasper Roe, Binh H. Vu, Darius Postma, Don Hickerson, James McGaughran, and Huy Q. Khuat. 2024. Simple techniques to bypass GenAI text detectors: implications for inclusive education. *International Journal of Educational Technology in Higher Education*, 21(1):53.

Alex Reinhart, Ben Markey, Michael Laudenbach, Kachatad Pantusen, Ronald Yurko, Gordon Weinberg, and David West Brown. 2025. Do LLMs write like humans? variation in grammatical and rhetorical styles. *Proceedings of the National Academy of Sciences*, 122(8):e2422455122.

Jasper Roe, Mike Perkins, Peter Bannister, Leon Furze, and James Wood. 2026. Dramaturgies of deception: AI humanizers and the performance of legitimacy in higher education assessment. *arXiv preprint arXiv:2605.02649*.

Jenna Russell, Marzena Karpinska, and Mohit Iyyer. 2025. People who frequently use ChatGPT for writing tasks are accurate and robust detectors of AI-generated text. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5342–5373, Vienna, Austria. Association for Computational Linguistics.

Vinu Sankar Sadasivan, Aounon Kumar, Sriram Balasubramanian, Wenxiao Wang, and Soheil Feizi. 2025. Can AI-generated text be reliably detected? stress testing AI text detectors under various attacks. *Transactions on Machine Learning Research*. ArXiv:2303.11156.

Chantal Shaib, Tuhin Chakrabarty, Diego Garcia-Olano, and Byron C. Wallace. 2026. Measuring AI "slop" in text. In *Third Conference on Language Modeling*. ArXiv:2509.19163.

Chantal Shaib, Yanai Elazar, Junyi Jessy Li, and Byron C. Wallace. 2024. Detection and measurement of syntactic templates in generated text. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 6416–6431, Miami, Florida, USA. Association for Computational Linguistics.

Zhengyang Shan, Yukyung Lee, and Sophie Hao. 2026. AI writers have a consistent stylometric footprint, but AI editors do not. *arXiv preprint arXiv:2608.27855*. ArXiv comment: EMNLP 2026 main conference; proceedings not yet published at time of access.

SuperAnnotate. 2024. Superannotate/ai-detector. Hugging Face model card, https://huggingface.co/SuperAnnotate/ai-detector. RoBERTa-large fine-tune; custom license; last updated 17 Dec 2024; accessed 2026-10-20.

Turnitin. 2026. How to access the AI writing report. Turnitin Guides, https://guides.turnitin.com/hc/en-us/articles/28457596598925-How-to-access-the-AI-Writing-Report. Last edited 6 March 2026; accessed 20 October 2026.

Peiyi Wang, Lei Li, Liang Chen, Zefan Cai, Dawei Zhu, Binghuai Lin, Yunbo Cao, Lingpeng Kong, Qi Liu, Tianyu Liu, and Zhifang Sui. 2024. Large language models are not fair evaluators. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9440–9450, Bangkok, Thailand. Association for Computational Linguistics.

Debora Weber-Wulff, Alla Anohina-Naumeca, Sonja Bjelobaba, Tomáš Foltýnek, Jean Guerrero-Dib, Olumide Popoola, Petr Šigut, and Lorna Waddington. 2023. Testing of detection tools for AI-generated text. *International Journal for Educational Integrity*, 19(1):26.

Wikipedia contributors. 2026. Wikipedia:signs of AI writing. https://en.wikipedia.org/wiki/Wikipedia:Signs_of_AI_writing. WikiProject AI Cleanup guide; accessed 2026-10-20.

Yale Undergraduate Admissions. 2026. Policy on the use of AI in application materials. https://admissions.yale.edu/ai-policy. Accessed 20 October 2026.

Peter Yang. 2026. no-ai-slop. GitHub repository, https://github.com/petergyang/no-ai-slop. Agent skill; MIT license; accessed 20 October 2026.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-judge with MT-bench and Chatbot Arena. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, volume 36, pages 46595–46623.

A Briefs

Set	Genre (n)	Topics
Development	Personal statement (4)	Four fictional students (Portuguese-American shoe-repair family; Nebraska church organist; Kazakh wild-apple collector; a Houston lion-dance performer), Common App prompts
	Short answer (4)	Same students, 250-word supplements (curiosity, roommate note, activity)
	Argument (4)	Phone bans; test retakes; group-project grades; part-time jobs
	Explainer (4)	mRNA vaccines; seasons; photosynthesis; sky colour and sunsets
	Literary analysis (4)	The Great Gatsby; Frankenstein; Macbeth; The Scarlet Letter (four verified quotes each)
Held-out	Personal statement (2)	Two new students: a Polish e-scooter repairer; an Arizona tamale-stand and weather-station student
	Short answer (2)	Why engineering; the world you come from
	Argument (2)	Later school start; graded homework
	Explainer (2)	Noise-cancelling headphones; leap years
	Literary analysis (2)	Romeo and Juliet; Pride and Prejudice
	Email (2)	Extension request to a teacher; cold email to a professor
	Cover letter (1)	Junior counselor at a library coding camp
	Blog post (1)	Robotics team at a regional competition
	Lab discussion (1)	Pendulum period against length
	News article (1)	Reusable cafeteria trays

Table 3: Writing briefs. Every brief names the only permitted source of specifics (interview notes, a facts file, four quotes, or none for arguments and explainers). All people are fictional. The held-out briefs were written after version 6 was fixed and share no example text with any version of the skill.